


Minimum Precision Requirements of Deep Neural Networks


Numerical Software Verification
July 20, 2020




Naresh R. Shanbhag

The Efficiency Challenge in Learning

Model Complexity



Data Movement Cost




fundamental question



how do we design **learning machines** that
operate at the limits of accuracy-robustness-
energy efficiency with **guarantees?**

Shanbhag Group Research Vectors

<http://shanbhag.ece.Illinois.edu>




low complexity algorithms



Machine Learning in Reduced Precision

MIT's Eyeriss



[ISSCC'16]

16b fixed-point
(inference)


Google's TPU



[ISCA'17]

8b fixed-point
(inference)
16b floating-point
(training)


Intel's NNP



[NIPS'17]

16b flexpoint
(training)

IBM's AI Core






[VLSI'18]

16b floating-point
(training)

Are these the minimum precisions required?

Can minimum precision requirements be determined analytically?

UIUC (Sakr, Shanbhag) - ICML 2017, ICASSP 2018, ICLR 2019, ICLR 2019 (with IBM)






- what are the minimum values of B_x^l , B_w^l , and B_y^l $\forall l$ such that the network accuracy is within a Δ of floating-point network accuracy?

Outline


- 1) “What are the minimum output precision requirements of a **dot product kernel** in order to meet a specific accuracy requirement at its output (**DP accuracy**)?”
- 2) “What are the minimum precision requirements of a **DNN** to meet a specific accuracy requirements at its output (**network accuracy**)?”
- 3) employ the above two insights to determine the precision limits of the recently proposed **in-memory computing (IMC) architectures**.

Minimum Output Precision Requirements for Dot Product



- what are the minimum values of B_x^l , B_w^l , and B_y^l $\forall l$ such that the network accuracy is within a Δ of floating-point network accuracy?

Why Output Precision B_y ?




$$y_q = \mathbf{w}^T \mathbf{x} + q_y$$


- B_y is the accumulator precision in digital architectures → accumulator complexity dominates power in low-precision DNNs
 - e.g., 32b accumulator 10× more power than a 3×1-b multiplier in 28nm CMOS – hence research on low-resolution accumulation [Sakr ICLR19; Wang NeurIPS’18]
- B_y is the ADC precision in in-memory architectures → ADCs can dominate (~80%) latency and power when implementing DNNs [Kim ISLPED’18, Rekhi DAC’20]

Quantization Noise Model

quantizer symbol



additive model




- additive model assumption: q_x is independent of x
- $SQNR$: signal-to-quantization noise ratio \rightarrow accuracy measure
- ζ : peak-to-average (power) ratio \rightarrow measure of ‘peakiness’ of signal distribution

$$SQNR_x = 10 \log_{10} \left[\frac{\sigma_x^2}{\sigma_{q_x}^2} \right]$$

$$SQNR_x(dB) = 6B_x + 4.78 - \zeta_x (dB)$$
$$\zeta_x = \frac{x_m}{\sigma_x}$$

Fixed-Point Dot Product



ideal FL output

y_o output quantization noise

$y_q = \mathbf{w}^T \mathbf{x} + q_{iy} + q_y$

input quantization noise
reflected at the output

$$y_q = \mathbf{w}^T \mathbf{x} + q_{iy} + q_y$$

Accuracy Metrics for Quantized Dot Products

$$SQNR_T = \left[\frac{1}{SQNR_{q_{iy}}} + \frac{1}{SQNR_y} \right]^{-1}$$


Limited by $SQNR_{q_{iy}}$

- Choose $SQNR_y(dB) \geq SQNR_{q_{iy}}(dB) + 9$ to minimize ($< 0.5dB$) its impact on $SQNR_T$

$$SQNR_y(dB) = 6B_y + 4.8 - [\zeta_x(dB) + \zeta_w(dB)] - 10 \log_{10}(N)$$

- But for fixed B_y : $SQNR_y(dB)$ **reduces with N** (N in hundreds in DNNs) → increase B_y
- But large B_y → leads to very large accumulator bit widths
- How to choose output precision B_y ?**

Bit Growth Criterion (BGC) for Choosing B_y





$$B_y = B_x + B_w + \log_2(N)$$

$$SQNR_y^{BGC}(dB) = 6(B_x + B_w) + 4.8 - [\zeta_x(dB) + \zeta_w(dB)] + 10 \log_{10}(N)$$

- commonly employed in digital architectures and network design
- B_y (accumulator precision) and $SQNR_y$ both **increase with N**


Proposed - Minimum Precision Criterion (MPC)



$$SQNR_y^{MPC}(dB) = 6B_y + 4.8 - \zeta_y^{MPC}(dB) - 10 \log_{10} \left(1 + p_c \frac{\sigma_{cc}^2}{\sigma_{qy}^2} \right)$$




- allow for a non-zero but small probability of clipping (p_c) – BGC avoids clipping
- exploits reduction in $\frac{\sigma}{\mu}$ of y_o with N (Central Limit Theorem)
- exhibits a trade-off between clipping noise and quantization noise

Comparing MPC and BGC



- MPC achieves the desired $SQNR_y^*$ with minimum precision ($B_y = 8$)
- BGC is a huge overkill → leads to very large accumulator bit widths ($B_y = 16$ to 20)
- tBGC (truncated BGC) needs $B_y = 12$ (still significant)
- Use MPC to assign minimum output (accumulator) precision

Input Precision Requirements for DNNs



- what are the minimum values of B_x^l , B_w^l , and B_y^l $\forall l$ such that the network accuracy is within a Δ of floating-point network accuracy?

Related Works


- much work on reduced precision machine learning since 2015 [stochastic rounding, BinaryNets, TernGrad, pruning...]

But.....


- largely based on heuristics – relying on the benevolence of SGD
- lacking theoretical guarantees on accuracy – try and hope it works!
- difficult to realize in H/W – complex and irregular arithmetic

Deep Learning in Finite Precision


Fixed-point **inference** with theoretical guarantees



Fixed-point **training** with close-to-minimal precisions



Fl.-pt. training with accumulation bit-width scaling




Sakr, Kim, Shanbhag
ICML 2017

Sakr & Shanbhag
ICASSP 2018

Sakr & Shanbhag
ICLR 2019

Sakr & Shanbhag
(with K. Gopalakrishnan [IBM])
ICLR 2019

Precision Analysis Framework



- no retraining; per-layer precision; activation vs. weight trade-off
- noise gains computed via one standard backprop iteration
- minimizing p_m done via **noise equalization (NE)**


Lesson 1 – Precision Trade-offs Captured Analytically

CIFAR-10 using ResNet 18

$$B_{A,l} = \log_2 \sqrt{\frac{E_{A,l}}{E_{\min}}} + B_{\min}$$

&

$$B_{W,l} = \log_2 \sqrt{\frac{E_{W,l}}{E_{\min}}} + B_{\min}$$



$$B_{A,l} = B_{W,l} = B_{\min} \text{ for } l = 1 \dots L$$



- bound **predicts** minimum precision within **1~2 bits**
- bound reveals **trade-offs** between network precisions
- trade-offs captured by **relative values of noise gains**

Lesson 2 – Precision Decreases with Depth

- weights typically require **more precision** than activations
- precision decreases because **early perturbations** are most **destructive**


CIFAR-10 using VGG-9

■ $E_{A,l}$; ■ $E_{W,l}$; ■ $B_{A,l}$; ■ $B_{W,l}$




Lesson 3 – BinaryNets are More Complex than Minimum Precision Networks

CIFAR-10 using VGG-9





- up to **3.5x lower** complexity & **2X** lower storage over BinaryNet at **iso-accuracy**
- empirically observed by [Moons, Verhelst]

Challenges in Fixed-point Training





- multiple forward quantization noise sources
- unknown gradient dynamic range
- instability due to quantization noise bias in updates
- back-propagation of quantization noise in activation gradients
- risk of premature stoppage of convergence



activations & weights spatio-temporally varying distributions



gradients spatio-temporally varying dynamic range



gradients & weights huge dynamic range mismatch



Criterion 1: equalization of quantization noise gains

$$B_{W_l} = \text{rnd} \left(\log_2 \left(\sqrt{\frac{E_{W_l \rightarrow p_m}}{E^{(\min)}}} \right) \right) + B^{(\min)}$$

$$B_{A_l} = \text{rnd} \left(\log_2 \left(\sqrt{\frac{E_{A_l \rightarrow p_m}}{E^{(\min)}}} \right) \right) + B^{(\min)}$$

Criterion 2: proper gradient clipping

$$r_{G_l^{(W)}} \geq 2\sigma_{G_l^{(W)}}^{(\max)}$$

$$r_{G_{l+1}^{(A)}} \geq 4\sigma_{G_{l+1}^{(A)}}^{(\max)}$$

Criterion 3: quantization bias elimination

$$\Delta_{G_l^{(W)}} < \frac{\sigma_{G_l^{(W)}}^{(\min)}}{4}$$

Criterion 4: back-propagated noise bound

$$\Delta_{G_{l+1}^{(A)}} < \frac{\Delta_{G_l^{(W)}}}{\sqrt{\lambda_{G_{l+1}^{(A)} \rightarrow G_l^{(W)}}^{(\max)}}} \left(\frac{|G_l^{(W)}|}{|G_{l+1}^{(A)}|} \right)^{1/4}$$

Criterion 5: accumulator stopping condition



$$r_{W_l^{(acc)}} \geq 2^{-B_{W_l}}$$

$$\Delta_{W_l^{(acc)}} < \gamma^{(\min)} \Delta_{G_l^{(W)}}$$

The Solution



- analytically guaranteed to provide close-to-minimal precision @ iso-accuracy
- requires a floating-point network to be probed during training

FX Training Converges with Close-to-Minimal Precisions



- FX training was believed to be **hard** due to **dynamic range issues** [Koester, NeurIPS'2017]
- proposed **FX training** is able to match **FL** training accuracy
- precision assignment found to be **nearly minimal**

Per-layer Precision Trends Vary



B_{W_j}
 B_{A_j}
 $B_{G_j^{(W)}}$
 $B_{G_j^{(A)}}$
 $B_{W_j^{acc}}$

- weight precision decreases from network input to output
- precisions of activation gradients and weight accumulators increase
- ResNets have uniform precision requirements per tensor type

Comparison w.r.t. Hyper-precision Reduction Techniques

	\mathcal{C}_W (10^6 b)	\mathcal{C}_A (10^6 b)	\mathcal{C}_M (10^9 FA)	\mathcal{C}_C (10^6 b)	Test Error		\mathcal{C}_W (10^6 b)	\mathcal{C}_A (10^6 b)	\mathcal{C}_M (10^9 FA)	\mathcal{C}_C (10^6 b)	Test Error
CIFAR-10 ConvNet						CIFAR-100 ResNet					
FL	148	9.3	94.4	49	12.02%	1789	97	4319	597	28.06%	
FX (C_o)	56.5	1.7	11.9	14	12.58%	750	25	776	216	27.43%	
BN	100	4.7	2.8	49	18.50%	1211	50	128	597	29.35%	
SQ	78.8	1.7	11.9	14	11.32%	1081	25	776	216	28.03%	
TG	102	9.3	94.4	3.1	12.49%	1230	97	4319	37.3	30.62%	


- feedforward binarization (BN) and gradient ternarization (TG) fail to match FL accuracy for same topology
- stochastic quantization (SQ) provides marginal returns
- BN, TG, SQ do not address the fundamental problem of realizing true FX training

Precision Requirements for In-Memory Architectures

UIUC 6T SRAM Deep In-memory IC Prototypes

100X EDP reduction over von Neumann equivalent* @ iso-accuracy


concept




[ICASSP'14]
(with Micron)

8b compute; 16kB SRAM in 65nm CMOS (UIUC)


multi-functional



random forest



on-chip learning



RNN attention ntwk



[arxiv'16, JSSC'18]

[ESSCIRC'17
JSSC'18]

[ISSCC'18,
JSSC'18]

[CICC'20]

To Speed Up AI, Mix Memory and Processing

New computing architectures aim to extend artificial intelligence from the cloud to smartphones

By Katherine Bourzac




Image: Sujan Gonugondla

Tearing Down Walls: This prototype features a new chip design called deep in-memory architecture.

State Circuits Conference (ISSCC), in San Francisco, he and others made their case for a new architecture that brings computing and memory closer together. The idea is not to replace the processor altogether but to add new functions to the memory that will make devices smarter without requiring more power.

ILLINOIS


Electrical & Computer Engineering

COLLEGE OF ENGINEERING

The Deep In-memory Architecture (DIMA)

Mingu Kang · Sujan Gonugondla · Naresh Shanbhag
Deep In-memory Architectures for Machine Learning

This book describes the recent innovation of deep in-memory architectures for realizing AI systems that operate at the edge of energy-latency-accuracy trade-offs. From first principles to lab prototypes, this book provides a comprehensive view of this emerging topic for both the practicing engineer in industry and the researcher in academia. The book is a journey into the exciting world of AI systems in hardware.




<https://spectrum.ieee.org/computing/hardware/to-speed-up-ai-mix-memory-and-processing>


Naresh Shanbhag – University of Illinois at Urbana-Champaign

Springer


In-memory ICs for Machine Learning is Hot!




[VLSI'16]




[JSSC'18]




[ISSCC'18]




[ISSCC'18]




[ISSCC'18]



[VLSI '19]




[ISSCC'20]




[ESSCIRC'17]




[ISSCC'18]




[arxiv'19]




[ISSCC'19]



[VLSI'19]




[ISSCC'20]




[ISSCC'20]



[VLSI'18]




[VLSI'18]




[VLSI '19]




[ISSCC '19]



[VLSI '19]





[ISSCC'20]




[CICC'20]

Research Questions

- Is there a common basis for these architectures?
- **What are their precision (compute SNR) limits? (BIG? for IMCs today)**




Fixed-Point Dot Product on IMCs



$$y = \mathbf{w}^T \mathbf{x} + q_i y + \eta_e + \eta_c + q_y$$


Annotations for the equation:

- ideal FL output:** A bracket above the term $\mathbf{w}^T \mathbf{x}$.
- clipping due to headroom limitations:** A bracket below the term $q_i y$, with a small waveform diagram above it.
- spatio-temporal analog noise sources:** A bracket below the noise terms $\eta_e + \eta_c + q_y$, with a small waveform diagram above it.



$$SNR_a = \frac{\sigma_{y_o}^2}{\sigma_{\eta_a}^2}$$

(analog SNR)




$$SNR_d = \frac{\sigma_{y_o}^2}{\sigma_{\eta_a}^2}$$

(digitization SNR)

$$SNR_T = \left[\frac{1}{SNR_a} + \frac{1}{SNR_d} \right]^{-1}$$

Limited by SNR_a

SNR Tradeoffs in Charge Summing Architectures




$$\text{SNR}_a = \frac{12\sigma_w^2 E[x^2]}{E[x^2]\Delta_w^2 + \sigma_w^2\Delta_x^2 + E[x^2]\frac{D_c}{4}\frac{\sigma_I^2}{I_o^2} + \sigma_w^2\frac{D_c}{4}\frac{\sigma_T^2}{T_o^2} + \frac{1}{N}\sigma_{\eta_c}^2}$$

- discharge current and pulse width trades-off with clipping noise
- clipping noise dominates as dimensionality N increases



SNR trades off with N and V_{WL}

ADC Precision Requirements



- precision limited by SNR_a

$$\text{SNR}_a(\text{dB}) - \text{SNR}_T(\text{dB}) \leq 0.5 \text{ dB}$$




$$B_y > \min \left(\log_2(k_{\text{clip}}), \frac{\text{SNR}_a(\text{dB}) + 16.6}{6} \right)$$


Summary

- Design of DNNs need not be trial-&-error based - analytical methods exist (for precision assignment) or can be developed
- use MPC & noise gain analysis to determine minimum precisions of DNNs
- parallel considerations for in-memory architectures – interplay between analog and quantization noise sources
- Next: design optimization framework? network accuracy vs. energy vs. latency vs....

Major challenge – engineered design of AI systems →
composability, interpretability, robustness, security, ethics, with guarantees



Machine Learning with Guarantees



Thank You!

<http://shanbhag.ece.uiuc.edu>